

CSE 493S Empirical Methods Assignment

Aadyant Maity and Rahul Bonthu

May 1, 2026

Part 0: implementation

Part 0 is basically the codebase the rest of the assignment builds on and everything lives under `deliverables/part_0/`: the usual scripts (`train.py`, `model.py`, `inference.py`, `tokenizer_utils.py`, `arith_data.py`) plus `part_0_1_contract.py` for the autograder-facing API we sanity-check with the quick configs under `deliverables/part_0/configs/`, for example:

```
python train.py --config configs/sanity_quick.yaml
python train.py --config configs/sanity_suffix_quick.yaml
```

(run from the `EmpiricalHW/` root, or point paths appropriately). Nothing fancy to report here besides “it trains and the contract loads checkpoints”; the interesting experiments start in Part 1.

Part 1.1: modular arithmetic datasets

For Part 1.1 we generated text files of equations $a \text{ op } b = c$ with arithmetic mod p . Addition and subtraction iterate all pairs (a, b) with $0 \leq a, b < p$. Division skips $b = 0$ and uses $b \in \{1, \dots, p-1\}$, so division sets are a bit smaller than add/sub for the same p . We shuffled once with seed 0 then split each file **80% / 10% / 10%** train, validation, and test in that order. Every `meta.json` matches the lines in the matching `.txt` files.

Split sizes

These counts match what is on disk under `deliverables/part_1/1.1/data/`.

Table 1: Number of lines per split for each operator and modulus.

p	op	train	val	test
97	+ / -	7527	940	942
97	/	7449	931	932
113	+ / -	10215	1276	1278
113	/	10124	1265	1267

To regenerate from the course workflow:

```
python arith_data.py --seed 0
```

Part 1.2: addition mod 97 (one-layer, deliverable run)

Part 1.2 is the addition-mod-97 setup with a one-layer model. Our submitted bundle lives under `deliverables/part_1/1.2/checkpoints/p97_add_1L_restart1_seed101/`, with final weights `ckpt_100000.pt` (`max_steps = 100000`, `save_every = 10000`), plus `tokenizer.json`, `metrics.csv`, and resolved config metadata.

Training curves (seed 101)

Figure 1 shows train/val/test loss, token accuracy, and equation accuracy for steps 0–100000 only (we pass `--max-step 100000` to `plot_metrics.py` so the figure matches the deliverable budget exactly).




Figure 1: Part 1.2 deliverable run: $p = 97$, addition, one layer, seed 101, 100k steps.

Numbers at the last logged step (seed 101)

At step 100000 (last row of the deliverable `metrics.csv`), we see roughly:

- Train loss ≈ 0.0136 , train token acc ≈ 0.996 .
- Val loss ≈ 0.0218 , val token acc ≈ 0.993 .
- Test loss ≈ 0.0275 , test token acc ≈ 0.993 .
- Train / val / test equation acc ≈ 0.985 , 0.974 , and 0.972 .

Part 1.3: division mod 97 (grokking-style run)

Part 1.3 trains on division modulo 97. The deliverable bundle is under `deliverables/part_1/1.3/checkpoints/part1_3_div_p97/`, with final weights `ckpt_300000.pt`, `tokenizer.json`, `metrics.csv`, and resolved config metadata.

Training curves

Figure 2 shows loss and accuracy vs. step from `metrics.csv`. We generate the PNG with `plot_metrics.py` from that run's CSV (repo-root script, `--max-step` optional).




Figure 2: Training / validation / test curves for the division mod 97 experiment (`part1_3_div_p97`).

Numbers at the last logged step

At step 300000 (last row of `metrics.csv`), the CSV reports roughly:

- Train loss ≈ 0.331 , train token acc ≈ 0.888 .
- Val loss ≈ 2.38 , val token acc ≈ 0.560 .
- Test loss ≈ 2.31 , test token acc ≈ 0.566 .
- Train equation acc ≈ 0.618 ; val/test equation acc stayed low on the last rows.

We are not over-interpreting these here; the figure is the main takeaway for us until we write up discussion elsewhere.

How to run inference

From `deliverables/part_1/1.3/code/`, we can use either free-form generation or `predict_answer` on integers. Example:

```
python inference.py --checkpoint_dir ../checkpoints/part1_3_div_p97 \  
    --prompt "3 / 5 = "
```

Part 1.4: ablations (division mod 97)

Part 1.4 compares two tweaks to the Part 1.3-style division setup ($p = 97$, one layer, 300k steps, seed 42). Each ablation has its own checkpoint folder under `deliverables/part_1/1.4/checkpoints/`.

Dropout (dropout = 0.2, weight_decay = 1.0)

Bundle: `part_1/1.4/checkpoints/part1_4_ablation_dropout/`.

Figure 3: `plot_metrics.py` with `--max-step 300000`.



Figure 3: Part 1.4 dropout ablation: division mod 97, dropout = 0.2, 300k steps.

At step 300000 (last row of that run's `metrics.csv`):

- Train loss ≈ 1.16 , train token acc ≈ 0.546 .
- Val loss ≈ 1.16 , val token acc ≈ 0.540 .

- Test loss ≈ 1.16 , test token acc ≈ 0.548 .
- Train / val / test equation acc ≈ 0.021 , 0.016 , and 0.021 .

Low weight decay (`weight_decay = 0.1`, `dropout = 0.0`)

Bundle: `part_1/1.4/checkpoints/part1_4_ablation_wd_low/`.

Figure 4: same `plot_metrics.py` setup with that run's `metrics.csv`.




Figure 4: Part 1.4 low weight decay ablation: division mod 97, `weight_decay = 0.1`, 300k steps.

At step 300000 (last row of that run's `metrics.csv`):

- Train loss ≈ 0.065 , train token acc ≈ 0.979 .
- Val loss ≈ 5.63 , val token acc ≈ 0.546 .
- Test loss ≈ 5.58 , test token acc ≈ 0.541 .
- Train equation acc ≈ 0.921 ; val / test equation acc ≈ 0.025 and 0.017 .

Part 2.1: Test-Time Scaling Warm-Up

For the warm-up in Part 2.1, we evaluated the Qwen3-4B model on the AIME 2024 dataset (30 problems) using greedy decoding. We measured the zero-shot accuracy under two conditions: allowing the model to use its thinking trace and suppressing the thinking trace (forcing an immediate answer).

Greedy Accuracy Baseline

As shown in Table 2, the model achieved 0% accuracy across both exact match and flexible extraction modes for all 30 problems.

Table 2: Greedy decoding accuracy on AIME 2024 (n=30).

Condition	Exact Match Accuracy	Flexible Extract Accuracy
With Thinking	0.000	0.000
Without Thinking	0.000	0.000

The uniform failure at zero accuracy is a direct result of the model’s behavior during greedy decoding (*do_sample = False*) specifically we observed a phenomenon of “mode collapse” or token attraction: once the thinking trace concluded, the model frequently fixated on a single repetitive integer (e.g., “78” or “108”) across many different prompts without the variance provided by sampling the model was unable to escape these local minima to arrive at the correct gold answers for the highly complex AIME problems.

Thinking Length Distribution

While accuracy was poor, the model successfully generated varied reasoning traces when thinking was enabled. We parsed these traces to measure the distribution of thinking tokens. Figure 5 illustrates the spread of reasoning compute used by the model before emitting a final answer.




Figure 5: Distribution of thinking trace lengths for Qwen3-4B on AIME 2024 problems.

The distribution characteristics are summarized below:

- **Mean:** $\approx 5,844$ tokens
- **Median:** $\approx 6,333$ tokens
- **Minimum:** 0 tokens (for problems where the model bypassed thinking)
- **Maximum:** 8,032 tokens

The varied lengths suggest that the model’s internal reasoning process scales dynamically with problem complexity, even when the final output fails to reach the correct answer. This baseline serves as the starting point for our subsequent scaling experiments in Part 2.2.

Part 3: Synthesis and analysis

When does each strategy help?

Train time is most informative when the data are *fixed, symbolic, and repeatable*. Part 1.1 supplies that setting: full modular splits for $p \in \{97, 113\}$ (e.g. 7527 train lines for $p = 97$ addition; division splits are slightly smaller because $b = 0$ is excluded). Part 1.2 shows train-time compute paying off on addition mod 97: at 100000 steps, train and test equation accuracy are both about 0.97–0.99, so generalization tracks training and division is harder. Part 1.3 at 300000 steps: train equation accuracy ≈ 0.62 while val and test equation accuracy remain near zero and val/test loss rise. Training still improves the training objective, but it does not produce held-out equation correctness within our budget. Part 1.4 isolates hyperparameters: dropout 0.2 yields near-chance token accuracy (≈ 0.54) and equation accuracy of order 10^{-2} across splits; weight decay 0.1 with no dropout drives train equation accuracy to ≈ 0.92 but leaves val and test equation accuracy at a few percent with val/test loss near 5.6. Train time then helps **memorization** or **stalled optimization**, depending on the knob, not reliably **test equation success**.

Test time in Part 2.1 (greedy Qwen3-4B on AIME 2024) behaves differently. Accuracy is 0% with and without thinking under exact and flexible scoring (Table 2), yet thinking traces are long on average (mean ≈ 5844 tokens, median ≈ 6333 , max 8032; Figure 5). Additional reasoning tokens under greedy decoding therefore did not improve measured accuracy here; they only increased internal computation. Test-time expansion is useful when decoding can still reach a correct, extractable answer. Part 2.1 also notes repetitive integer outputs under greedy settings, which is consistent with compute being spent without improving the graded output.

Summary. Train-time steps aligned well with generalization for Part 1.2; for Parts 1.3–1.4 the same budget tells a split or stagnant test story. Inference-time length in Part 2.1 did not translate into accuracy gains under our greedy protocol.

Failure modes and what they share

We highlight three patterns visible in this report: (1) Part 1.3 and the low-WD Part 1.4 run show strong train-side progress with weak val/test equation accuracy; (2) Part 2.1 shows long thinking traces with 0% greedy accuracy when thinking is enabled; (3) Part 2.1 discusses repetitive final-token behavior under greedy decoding. We do not include Part 2.2 budgets in this PDF, so we do not empirically demonstrate accuracy *dropping* as a function of longer thinking here; we only observe that longer traces did not help under greedy decoding.

The common structure is **mismatch between compute and the evaluation functional** optimization reduces training loss without yielding correct full equations on held-out modular lines (Parts 1.3–1.4) generation increases thinking tokens without passing the AIME extractors (Part 2.1) the assignment also notes that parallel samples can lack diversity we have not plotted parallel runs in this document conceptually duplicated wrong samples resemble repeated wrong greedy strings more draws need not increase coverage of the correct answer set.

Taken together, compute and generalization are **not** linked by a fixed monotone map. Either more steps or more tokens can help when they preserve a path to solutions that remain valid under the test metric; otherwise they mainly deepen an unhelpful attractor (memorization without equation generalization, or repetitive incorrect completions).

Budget allocation across difficulties

Consider a single pretrained reasoning model a math benchmark with easy medium and hard items and a fixed total compute budget. **Easy items:** prioritize inexpensive inference (short or no thinking, greedy or low temperature). Reserve finetuning only if errors are systematic format issues rather than reasoning depth. **Medium items:** allocate a moderate sequential thinking budget first informed by Part 2.1’s observation that traces can be long even when wrong; hold a smaller reserve for targeted continued pretraining or finetuning if diagnostics reveal gaps resembling modular division (Parts 1.3–1.4). Monitor val and test curves as in Part 1.2 and avoid configurations like Part 1.4 dropout that consume steps without progress. **Hard items:** emphasize search over naive length (parallel samples, temperature, or alternative decoders) since Part 2.1 shows greedy length alone did not lift accuracy. Train-time spend is secondary unless auxiliary structured tasks (Parts 1.1–1.2) are shown to transfer; Parts 1.3–1.4 caution that division-style structure remains difficult.

Oracle for partial quality: if partial solutions could be scored, we would reallocate inference compute from blind restarts toward **verification and revision** of candidate answers. Parts 1.3–1.4 already show that train loss can improve while equation-level correctness lags, so an oracle would help tie extra tokens to checks that correlate with final correctness rather than to train loss alone.

This allocation is one reasonable policy consistent with Parts 1.1–1.4 and 2.1; it is not claimed to be uniquely optimal.